

The Retention Engine — Spaced Repetition Over the Manual's Own Facts

Chapter F11.7 — 7,068 Cards, and the Schedule That Makes Them Worth Having

Module	Foundation Layer — Module 11 — PYQ Analysis, Mocks and Final Revision
Chapter	Chapter F11.7
Topic	Retention scheduling, card discipline and the decay problem
Version	v1.0
Compiled	22 August 2026
Exam stage	Recruitment Test (RT)

MUST MASTER · THE ENEMY IS NO LONGER COVERAGE

This manual now holds around **230 hours of reading** and **1,459 practice questions**. F11.3 §1 already names the consequence, and Register §5 records it: with the syllabus covered, **the binding constraint has moved from coverage to decay**. The risk is not that you will fail to reach a module. It is that the module you read in September will have gone by November.

Nothing in the manual fought that. Wall sheets are the right *content* for the fight and the wrong *mechanism* — a wall sheet is re-read, and re-reading is the least efficient form of study known, because recognising a fact feels exactly like knowing it.

This chapter supplies the mechanism. decks/ holds **7,068 cards** generated from the manual's own sources — **all twelve modules**, since the Module 1, 2, 3 and 5 sources arrived. §2 is why they are shaped as they are, §4 is the schedule, and **§5 is the part most people skip and should not: what to do when a card fails**.

Read §4.3 before you start. 7,068 cards is a library, not a target. Which slice you run is the single decision that determines whether this works.

Table of Contents

0. How This Chapter Works

1. Why Re-reading Fails and Retrieval Does Not

2. How the Cards Were Built

2.1 The seven decks

2.2 Modules 1, 2, 3 and 5 — present, but as PDFs

2.3 The extra gate on PDF-derived cards

3. Setting Up, Once

4. The Schedule

4.1 The daily habit

4.2 Which decks, when

4.3 Choosing scope — the decision that matters most

5. When a Card Fails

6. What This Does Not Do

7. Five-Minute Wall Sheet

0. How This Chapter Works

STUDY SESSION · CHAPTER F11.7

One session, forty minutes, done once. This is a setup chapter, not a study chapter. You read it, you import seven files, you choose your scope from §4.3, and you start a daily habit.

Then you never read it again except §5, which you will need in week three when the first cards start failing and you are tempted to do the wrong thing about it.

The decks are built by `build/decks.py` from `src/`. They are not a separate body of content to be maintained, and that is deliberate: **a fact that changes in a chapter changes in the deck on the next build**. Given how much of this syllabus moved in 2025–26, any deck that had to be hand-edited would be wrong within a month.

1. Why Re-reading Fails and Retrieval Does Not

The finding is old, replicated and inconvenient: **testing yourself on material produces far more durable memory than reviewing it**, and it does so while feeling *worse* at the time. Re-reading produces fluency — the text runs smoothly, nothing snags — and fluency is mistaken for mastery. Retrieval produces effort, hesitation and the occasional blank, all of which feel like failure and are in fact the mechanism working.

Two consequences for this exam specifically.

First, the wall sheets are being used wrongly by almost everyone. They are excellent — dense, atomic, already stripped to the examinable core. But reading one is recognition practice, and the paper does not test recognition. F11.4 §2 makes the point from the other direction: the whole trap taxonomy works by supplying options that *feel* familiar. **A candidate trained on recognition is trained on precisely the faculty the paper exploits.**

Second, spacing beats massing. Reviewing a fact on days 1, 3, 8 and 21 produces stronger retention than reviewing it four times on day 1, at identical total cost. This is why a daily twenty-minute habit outperforms a Sunday three-hour revision block, and why starting the habit late is expensive in a way that cannot be made up by working harder in the last fortnight.

THINK LIKE UPSC · WHY THIS PAPER PUNISHES DECAY UNUSUALLY HARD

Most exams let partial knowledge earn partial credit. This one does not. F11.1 §1.2 shows a wrong answer costs **0.8333** against a correct answer's **2.5** — so a half-remembered fact is not worth a third of a mark, it is worth **less than nothing** if it produces false confidence.

And F11.4 §3's T-09 shows the specific mechanism: the distractors are *superseded values* — ₹6,500 for ₹15,000, 1.61% for 1.16%. A decayed memory does not return "I don't know". **It returns the neighbouring value with the same feeling of familiarity.** That is why decay on this paper converts directly into negative marks, and why the numbers deck is reviewed on the shortest interval of the seven.

2. How the Cards Were Built

Four rules governed generation. They matter because they are the difference between a deck that trains the exam and a deck that trains a feeling of competence.

MUST MASTER · THE FOUR RULES

- 1. Recognition is not recall.** An MCQ's answer is on the page. **No card front reproduces options.** Every card either deletes the fact from its own sentence, or asks a question whose answer is not visible.
- 2. One fact to a card.** A wall-sheet line reading $\text{BRASS} = \text{Cu} + \text{Zn} \cdot \text{BRONZE} = \text{Cu} + \text{Sn}$ becomes two cards. A card holding two facts fails for the wrong reason and hides which half you lost.
- 3. Discrimination beats assertion.** Knowing ₹15,000 is nearly worthless while ₹6,500 feels equally familiar. So the confusion tables become their own deck, and those cards **name the distractor and require you to refuse it.**
- 4. Never fabricate a question.** Where a line had no clozeable target and no natural split it was **skipped and counted** — **3,268** lines in all, once the page furniture and the welded PDF lines of §2.2 are added — rather than turned into a vague prompt. The count prints on every build, so the gap stays visible instead of being papered over.

2.1 The seven decks

Deck	Cards	Built from	Why it exists
1 Numbers and Dates	2,994	Every figure in the wall sheets and explanations	Numbers decay fastest and T-09 attacks them hardest
2 Discrimination	131	The confusion and superseded-instrument tables	Highest value per card in the whole set
3 Core Facts	2,611	Wall sheets and one-minute revisions	The examinable core, one fact at a time
4 Question Recall	967	Practice items, converted to free recall	Turns 1,459 recognition items into retrieval
5 Definitions and Pairs	293	Term-to-meaning and List-I/List-II tables	Matching items are pure paired recall
6 Traps	30	T-01 to T-30 of F11.4	Technique decays too, and faster than facts
7 Interview	42	The one-line interview answers	Spoken fluency, prompted aloud

Deck 2 is small and you should not read that as unimportant. **131 cards covering the exact confusions this paper is built on is the highest-yield hour in the manual.**

2.2 Modules 1, 2, 3 and 5 — present, but as PDFs

Those 44 chapters — **97.5 of the 300 planning marks** — arrived as PDFs rather than markdown, and are read from `incoming/` by the same build. They contribute **2,643 cards**. What they do *not* contribute is the point.

Deck	From the eight markdown modules	From the four PDF modules
1 Numbers and Dates	1,150	1,844
3 Core Facts	1,812	799
2 Discrimination	131	0
4 Question Recall	967	0
5 Definitions and Pairs	293	0

A PDF has thrown away the structure the other decks are made of. Deck 2 needs a confusion table to still be a **table**; Deck 4 needs the `| Q | Ans | Explanation |` grid to survive as a grid. Flattened to text, both are unrecoverable — so those four modules yield facts and figures, and nothing else.

IMPORTANT · THE 52-MARK SUBJECTS GOT FACTS BUT NOT DISCRIMINATIONS

Labour law and social security are **52 marks** and the substance of the post, and they are exactly the subjects where F11.4's thesis bites hardest — ₹15,000 against ₹6,500, 1.16% against 1.61%, the 1952 Act against Chapter III of the Code. **Those are discriminations, and the PDF path produced none of them.**

You are not unprotected: Deck 2's 131 cards come from **F11.4 §5**, which was written for precisely these confusions and covers the provident fund, pension and insurance numbers directly. But it is a curated selection, not the whole of Modules 2 and 3.

If markdown for those four modules exists anywhere, supply it. One rebuild converts them from fact-and-figure sources into full ones. Measured on a single chapter built both ways: markdown 97 cards, its PDF 71.

2.3 The extra gate on PDF-derived cards

Two defects are specific to reading a PDF, and both were producing cards before they were caught.

Page furniture. A running head and footer appear on every page, and `Chapter F1.13 Page 25` mined as a fact. **393 lines** rejected.

Welded lines. A PDF has no paragraph structure, so wall-sheet entries that were visually separate arrive on one line: `1.1 Section 1(3) — the trigger 1.2 Sections 16 and 17 .` Clozing that produces a card holding three facts. Rejected on length, on a run of several distinct capitalised labels, or on an enumeration of figures — **897 lines** in all.

TRAP · A BAD CARD IS WORSE THAN NO CARD, SO THE GATE IS DELIBERATELY HARSH

Before the gate the PDF modules yielded **3,980** cards; after it, **2,643**. A third was thrown away, and that is the right trade.

A card you cannot answer because it holds three facts does not teach you the three facts. It teaches **hesitation**, it consumes a review slot every day, and — because you will eventually recognise its shape rather than recall its content — it manufactures exactly the false familiarity that F11.4 §3's distractors are built to exploit. **Sampling every deck by hand is therefore part of the build, not an optional check.**

3. Setting Up, Once

1. Install **Anki** — desktop for import, phone for review. The phone is where the twenty minutes actually get done, in queues and between things.
2. **File » Import**, select each of the seven files in `decks/`, click Import. No column mapping: the note type, deck and tag column are declared inside each file.
3. Set new cards per day to **40** and reviews to **unlimited**. Then read §4.2 for what that actually costs per day, and §4.3 for what to cut if the answer is more time than you have.
4. Turn **off** the option that buries siblings, so both cloze deletions of one fact can appear in a day.

CHECK YOURSELF · SETUP

Import is correct if the deck list shows seven subdecks under **APFC**, the numbers and core-facts decks show as **Cloze**, and a card from Deck 2 shows a tick line and a cross line on its back. If instead you see one deck called *Default* with 7,068 notes, the header directives were stripped — re-import the file unmodified.

4. The Schedule

4.1 The daily habit

Every day, on the phone. Not an hour on some days and none on others: the interval arithmetic assumes daily attendance, and a two-day gap does not cost two days of progress — it collapses the scheduling for every card that came due in the gap.

How long it costs is a function of scope, and the arithmetic is worth doing before you start. New cards generate reviews for weeks afterwards, at roughly eight to ten reviews per new card before an interval gets long. So the daily load settles at something like ten times your new-card rate:

New cards a day	Steady-state reviews	Minutes a day	Clears 7,068 in
20	~180	20–25	50 weeks
40	~350	35–45	25 weeks
60	~520	55–70	17 weeks

Nobody sustains the third row. If 25 weeks is more than you have — and for an exam this cycle it is — the answer is §4.3, not the bottom row.

Order matters. Do **due reviews first**, always, then new cards with whatever time is left. A backlog of reviews is decay in progress; new cards are only ever an addition.

4.2 Which decks, when

Phase	Decks running	New cards a day	Emphasis
Weeks 1–2	1, 2	40	Numbers and discriminations only. These are the cards whose absence costs marks
Weeks 3–4	1, 2, 6	40	Traps join — 30 cards, and they protect every other mark. This is Core
Weeks 5–10	+ 3, filtered	40	Deck 3 by tag, weakest modules first. Volume peaks here. Mock 2 at week 8
Weeks 11–12	+ 5, and 4 if room	25	Rate falls as the review load matures. Mock 3 at week 12
Last 2 weeks	reviews only	0	No new cards. Clearing due reviews only, which shortens every day
Interview gap	7	42 total	Deck 7 begins the day the RT ends, spoken aloud, not read

TRAP · DO NOT ADD NEW CARDS IN THE LAST TWO WEEKS

New cards generate reviews on a 1-, 3-, 8-day rhythm, so a card added fourteen days out is still consuming attention on the morning of the paper, and consuming it at the worst moment — while it is still unstable.

F11.3 §4 sets the last four weeks for retrieval, timing and error correction. **Adding cards inside two weeks is coverage work disguised as revision work.** Stop at day 14 and let the queue drain.

4.3 Choosing scope — the decision that matters most

Do not raise the daily new-card count. It inflates reviews for weeks afterwards and collapses in week three, usually taking the habit with it. Cut scope instead, in this order.

Tier	What you run	Cards	At 40/day
Core	Decks 1, 2, 6 whole	3,155	11 weeks
+ weak subjects	Core, plus Deck 3 filtered by tag to your four weakest modules	~4,500	16 weeks
+ definitions	The above, plus Deck 5	~4,800	17 weeks
Everything	All seven decks	7,068	25 weeks

Start at Core and add upward only if the daily load is comfortable after two weeks. Anki filters by tag, so `tag:syllabus::labour-law` or `tag:APFC::M03` narrows Deck 3 without deleting anything — the cards stay available if you later have room.

Drop Deck 4 first if you are short. Its 967 cards are the lowest-yield in the set for you specifically, because their content is reachable by sitting F11.2, F11.5 and F11.6 — which you are doing anyway, under timed conditions, which is strictly better practice than a flashcard.

MUST MASTER · WHY CORE IS DECKS 1, 2 AND 6 AND NOT THE BIGGEST ONES

Deck 3 is the largest at 2,611 cards and it is **not** in Core. That is deliberate.

Deck 1 is **numbers**, which decay fastest and which T-09 attacks by offering the superseded value. Deck 2 is **discriminations**, which is the failure mode named in F11.4 §2. Deck 6 is **technique**, 30 cards, which protects every other mark in the paper.

Deck 3 is the examinable core and it is genuinely valuable — but it is also the material most likely to be carried by a chapter you have read recently, and the least likely to produce a *confidently wrong* answer. **Core is chosen by what costs marks when it decays, not by what covers the most syllabus.**

5. When a Card Fails

This is the section that decides whether the engine works, and it is where the instinct is wrong.

The instinct is to re-read the chapter. Do not. A single lapse is normal and is the schedule working — the algorithm found the edge of your retention, which is exactly what it is for. Press *Again* and move on.

Act on the third lapse of the same card, not the first. Three failures is no longer decay, it is a defective card or a genuinely unlearned fact. Then, and only then, do one of these three things.

The card	The diagnosis	The action
Fails because two facts are tangled	You have a confusion , not a gap	Find its partner in Deck 2. If absent, write one card naming both, in the Deck 2 shape
Fails because the fact is isolated	No structure to hang it on	Read that chapter's §, once. This is the only case that justifies going back to prose
Fails because the card is bad	Ambiguous cloze, missing context	Fix or delete it. A bad card is worse than no card — it trains hesitation

THINK LIKE UPSC · A LAPSE IS INFORMATION, AND IT IS THE SAME INFORMATION F11.3 §4.2 WANTS

F11.4 §9.4 builds an error signature from mock answers; F11.3 §4.2 keeps a paper error log. **A repeated card lapse is the same signal arriving earlier and cheaper** — before it costs 0.8333 in a mock.

So read the two together. If Deck 2 keeps failing on superseded instruments and your mock errors cluster on **T-09**, that is not two problems. It is one, and it has a name, and F11.4 §5.1 is the fix.

6. What This Does Not Do

Stated plainly, because a deck invites over-confidence.

It does not teach. Every card was generated from a chapter you are expected to have read. A card is a retrieval handle on existing knowledge, and cards on an unread chapter produce memorised strings that collapse the moment the paper rephrases them.

It does not replace mocks. Cards train fact retrieval; the paper additionally demands 60-second pacing, OMR discipline, and the elimination judgment of F11.1 §1.2. F11.2, F11.5 and F11.6 remain the instrument for those. **A candidate with 7,068 mature cards and no mocks will still mismanage the two hours.**

It covers the 97.5-mark modules only as facts and figures. No discriminations, no question recall — see §2.2, and supply markdown if it exists.

And it does not fix a wrong fact. The cards inherit their sources exactly. Where this manual is out of date the deck will drill the error with perfect efficiency — which is why every figure carries its instrument and vintage, and why `decks/` is regenerated rather than hand-maintained.

7. Five-Minute Wall Sheet

FIVE-MINUTE WALL SHEET · CHAPTER F11.7

THE ENEMY IS DECAY, NOT COVERAGE

RE-READING BUILDS FLUENCY · RETRIEVAL BUILDS MEMORY · FLUENCY FEELS LIKE MASTERY AND IS NOT

SPACED BEATS MASSED AT IDENTICAL TOTAL COST — days 1, 3, 8, 21

DECAY HERE RETURNS THE NEIGHBOURING VALUE, NOT A BLANK — that is why it costs 0.8333

7,068 CARDS · SEVEN DECKS · ALL TWELVE MODULES · GENERATED, NEVER HAND-EDITED

DECK 1 NUMBERS SHORTEST INTERVAL · DECK 2 DISCRIMINATION HIGHEST YIELD PER CARD

TWENTY MINUTES DAILY · DUE REVIEWS FIRST, ALWAYS · NEW CARDS ONLY WITH TIME LEFT

40 NEW A DAY = ~350 REVIEWS, 35–45 MIN · CLEARS ALL 7,068 IN 25 WEEKS

CORE = DECKS 1, 2, 6 = 3,155 CARDS = 11 WEEKS · START THERE

SHORT OF TIME » CUT SCOPE, NEVER RAISE THE DAILY COUNT · DROP DECK 4 FIRST

NO NEW CARDS INSIDE THE LAST FOURTEEN DAYS

TRAPS DECK JOINS LATE — TECHNIQUE STABILISES AFTER FACTS DO

ACT ON THE THIRD LAPSE, NOT THE FIRST

THREE LAPSES = CONFUSION, ISOLATION OR A BAD CARD — DIAGNOSE BEFORE RE-READING

A BAD CARD IS WORSE THAN NO CARD

DECK 7 STARTS THE DAY THE RT ENDS · SPOKEN ALOUD, NOT READ

MODULES 1, 2, 3, 5 ARE PDFs = FACTS AND FIGURES ONLY, NO DISCRIMINATIONS · MARKDOWN 100%, PDF 73%

ONE LINE FOR THE INTERVIEW Asked how I prepared, I would say I stopped re-reading once I realised recognition was the exact faculty the paper is designed to exploit, and moved to daily retrieval over about four thousand single-fact cards — which had the useful side effect of showing me, in week three rather than in the hall, that my errors were clustered on superseded figures rather than spread across the syllabus.